

# **Database Systems**

**(CS 355 / CE 373)**

Dr. Umer Tariq  
Assistant Professor,  
Dhanani School of Science & Engineering,  
Habib University

# Acknowledgements

- Many slides have been borrowed from the official lecture slides accompanying the textbook:

Database System Concepts, (2019), Seventh Edition,

Avi Silberschatz, Henry F. Korth, S. Sudarshan

McGraw-Hill, ISBN 9780078022159

The original lecture slides are available at:

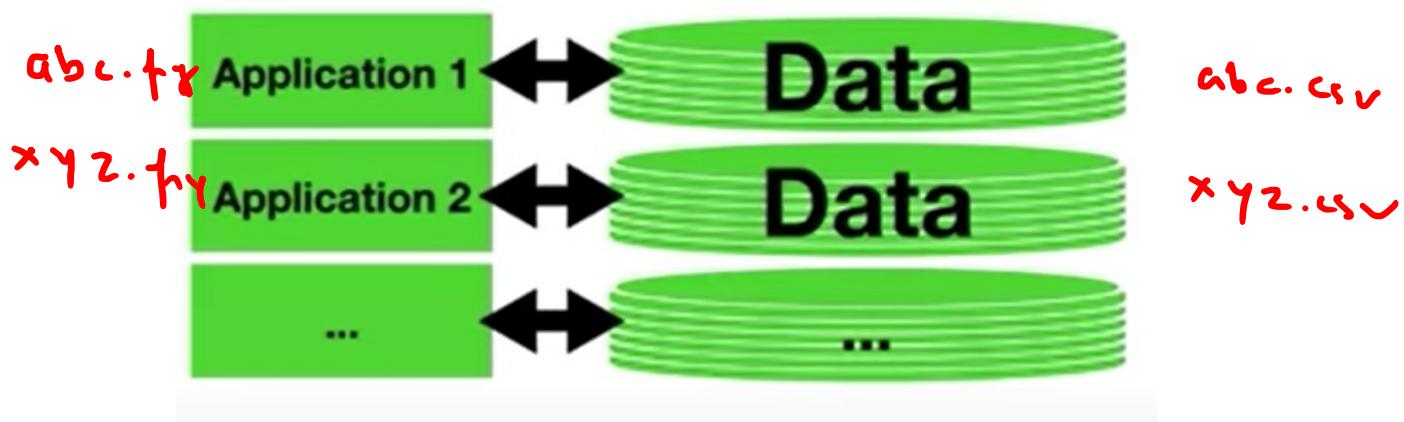
<https://www.db-book.com/>

- Some of the slides have been borrowed from the lectures by Dr. Immanuel Trummer (Cornell University). Available at: ([www.itrummer.org](http://www.itrummer.org))

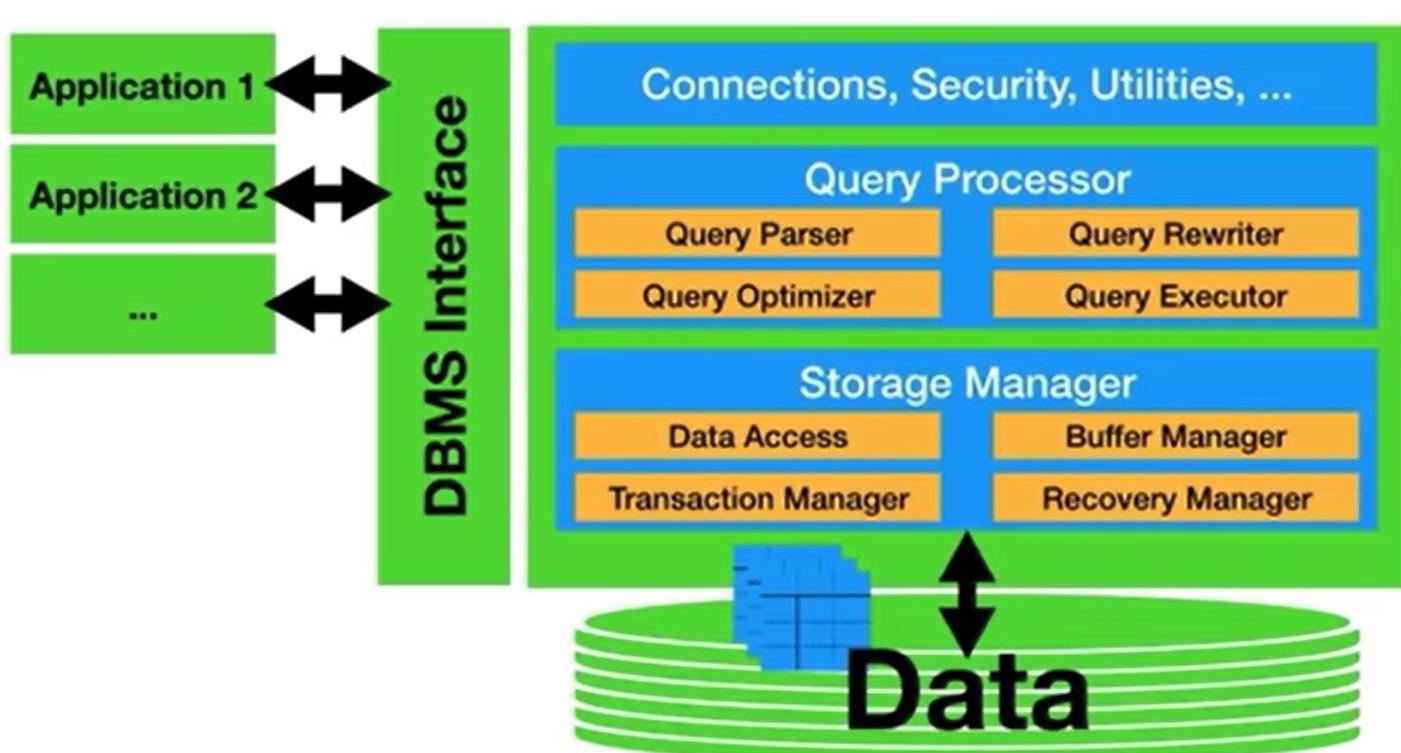
# Outline: Week 2

- Introduction to Relational Data Model
- Structure of Relational Databases
- Relational Database Schema
- Keys
- Schema Diagrams

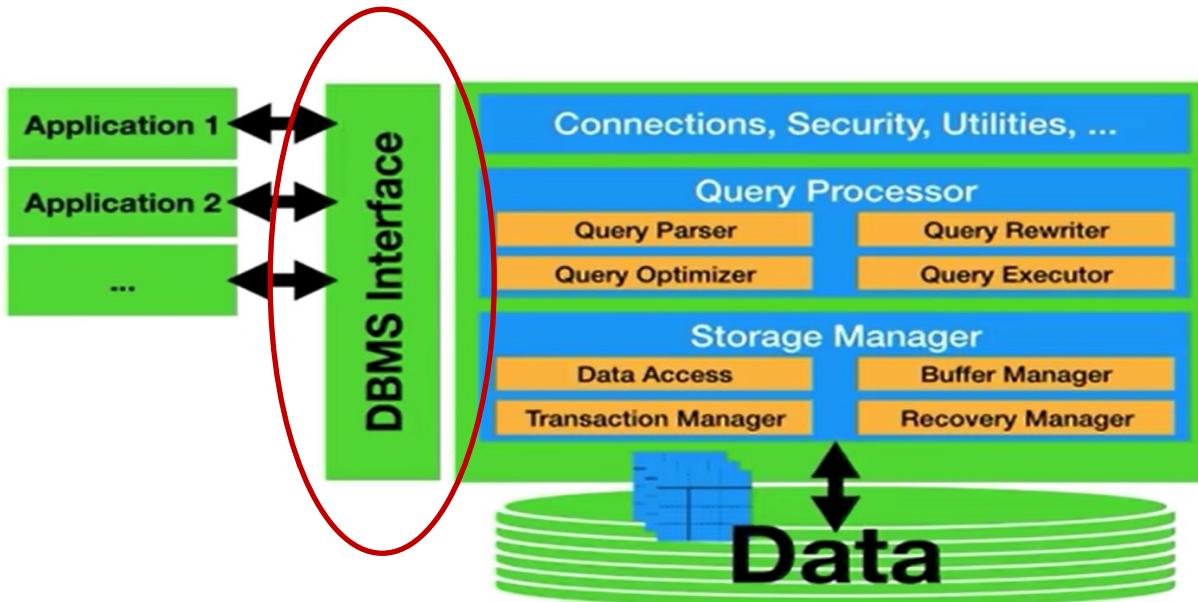
# File-Based Approach



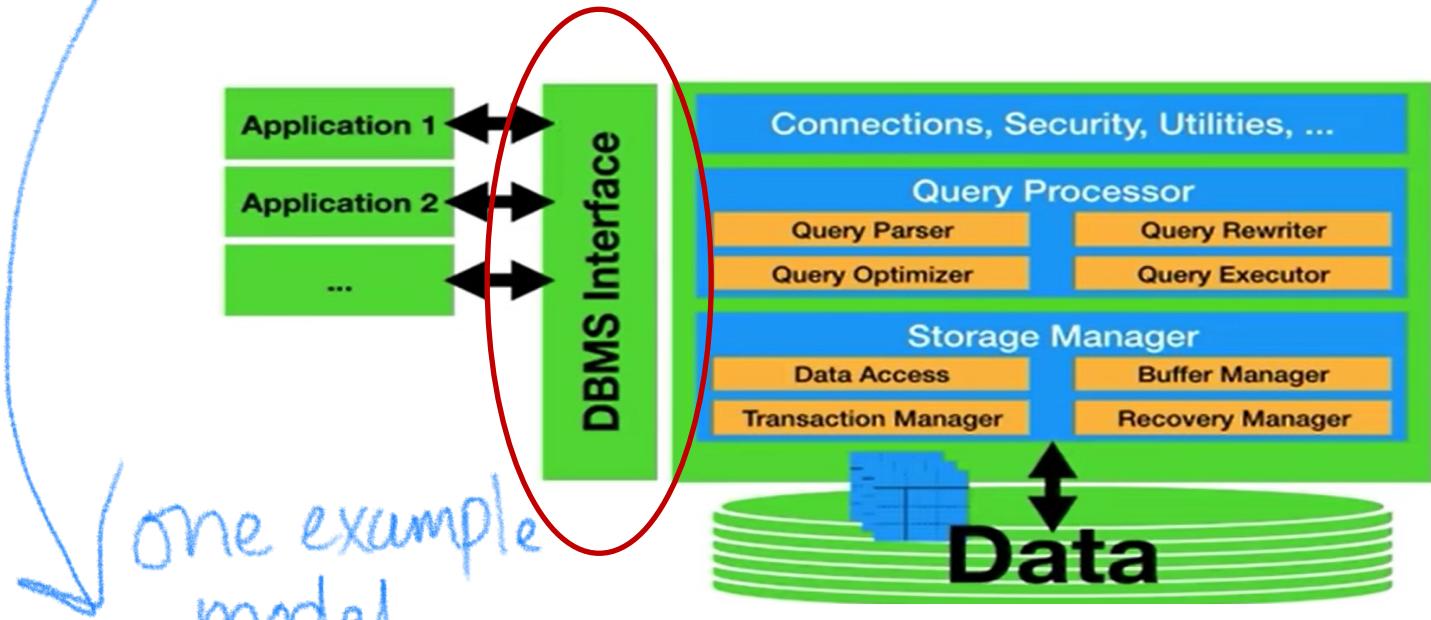
# Database Management System (DBMS)



# What should be the DBMS Interface?



## What should be the DBMS Interface?



- One example model
- "Data Model"
    - A collection of conceptual tools for describing data, data relationships, data semantics, and consistency constraints.

# The Relational Model

“TABLE - based  
Data Model”

- The relational model uses a collection of tables to represent both data and the relationships among those data.
- Its conceptual simplicity has led to its widespread adoption; a vast majority of database products are based on the relational model.

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

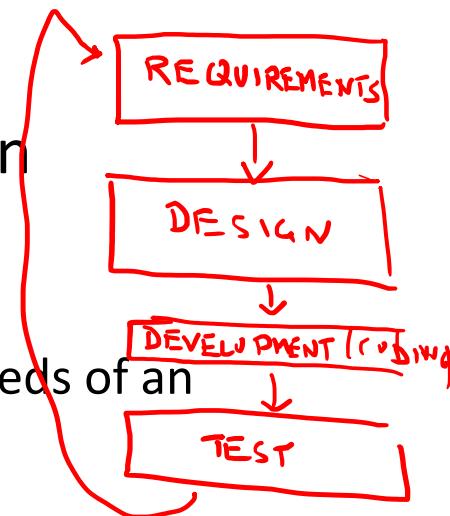
(a) The *instructor* table

<i>dept_name</i>	<i>building</i>	<i>budget</i>
Comp. Sci.	Taylor	100000
Biology	Watson	90000
Elec. Eng.	Taylor	85000
Music	Packard	80000
Finance	Painter	120000
History	Painter	50000
Physics	Watson	70000

(b) The *department* table

# Application Design vs Database Design

- Database-based applications are developed to meet the needs of an enterprise/business
- The design of a complete database application (so that it meets the needs of the business) requires attention to a broad set of issues such as:
  - What data to store
  - What should be the user interface
  - What should be the architecture of the application
- “Database Design”
  - Limited to design of database schema (the set of tables that will be used to store the data).



# The Relational Model

- The relational model uses a collection of tables to represent both data and the relationships among those data.

Table **instructor**

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

(83821, Brandt, Comp. Sci., 92000)

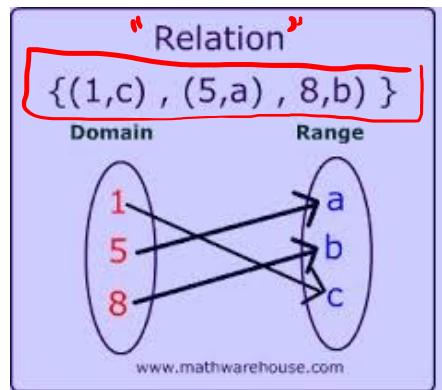
Figure 2.1 The *instructor* relation.

## Why is It Called the “Relational” Model?

- Let's talk some math!

$$\{1, 5, 8\} \quad \{a, b, c\}$$

$$\{(5, b), (1, a), (8, c)\}$$



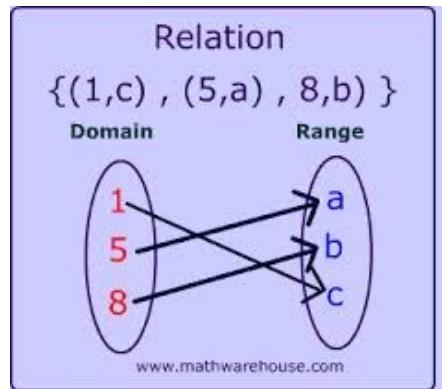
- The relation between two sets is a collection of pairs (2-tuples) containing one object from each set.

①

②

# Why is It Called the “Relational” Model?

- Let's talk some math!



$\{ \underbrace{(1,a,A), (5,b,B), \dots}_{\text{3-tuple}} \}$

$\circled{A} \cdot B \cdot C$

- The relation between two sets is a set of pairs (2-tuples) containing one object from each set.
- The relation between three sets is a set of 3-tuples containing one object from each set.
- The relation between  $n$  sets is a set of  $n$ -tuples containing one object from each set.

# Why is It Called the “Relational” Model?

`dept_name = { CompSci, Biology, EE, Music, ... }`

`building = { Taylor, Watson, Packard, Painter }`

`budget = { Set of All integer }`

`Relation  
on three sets = { (Comp.Sci, Taylor, 100000), (Biology, Watson, 90000),  
(EE, Taylor, 85000), ... }`

<code>dept_name</code>	<code>building</code>	<code>budget</code>
Comp. Sci.	Taylor	100000
Biology	Watson	90000
Elec. Eng.	Taylor	85000
Music	Packard	80000
Finance	Painter	120000
History	Painter	50000
Physics	Watson	70000

(b) The *department* table

- The relation is a set of tuples.

# The Relational Model

- The relational model uses a collection of tables to represent both data and the relationships among those data.

**Table/ Relation**  
instructor

The diagram illustrates the *instructor* relation as a table. A large curly brace on the left side groups the entire table under the heading "Table/ Relation" and "instructor". The table has five columns: *ID*, *name*, *dept\_name*, *salary*, and an unnamed column represented by a blue-bordered box. The *dept\_name* column is highlighted with a red border and labeled "Column / Attribute dept\_name". The row for the last entry (ID 83821) is also highlighted with a blue border and labeled "Row / Tuple (83821, Brandt, Comp. Sci., 92000)".

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>	
10101	Srinivasan	Comp. Sci.	65000	
12121	Wu	Finance	90000	
15151	Mozart	Music	40000	
22222	Einstein	Physics	95000	
32343	El Said	History	60000	
33456	Gold	Physics	87000	
45565	Katz	Comp. Sci.	75000	
58583	Califieri	History	62000	
76543	Singh	Finance	80000	
76766	Crick	Biology	72000	
83821	Brandt	Comp. Sci.	92000	
98345	Kim	Elec. Eng.	80000	

**Column / Attribute**  
*dept\_name*

**Row / Tuple**  
(83821, Brandt, Comp. Sci., 92000)

**Figure 2.1** The *instructor* relation.

# Let's Practice Our Terminology!

- Identify some attributes of the *course* relation?

*course\_id*, *title*

- Identify any tuple in the *course* relation?

(BIO-101, Intro.to Biology, Biology, 4)

- Are there any attributes that have a unique value in each tuple of this relation?

*course\_id*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-101	Intro. to Biology	Biology	4
BIO-301	Genetics	Biology	4
BIO-399	Computational Biology	Biology	3
CS-101	Intro. to Computer Science	Comp. Sci.	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3
CS-319	Image Processing	Comp. Sci.	3
CS-347	Database System Concepts	Comp. Sci.	3
EE-181	Intro. to Digital Systems	Elec. Eng.	3
FIN-201	Investment Banking	Finance	3
HIS-351	World History	History	3
MU-199	Music Video Production	Music	3
PHY-101	Physical Principles	Physics	4

Figure 2.2 The *course* relation.

Comp. ID	Comp. Name	Office Location
1	ENLRD	Lahore, Karachi, Islamabad
2	Unilever	Lahore, Karachi
-	-	-



## Properties of Attributes

Comp. ID	Comp. Name	Office Location
1	ENLRD	Lahore
1	ENLRD	Karachi
1	ENLRD	Islamabad
2	Unilever	Lahore
2	Unilever	Karachi

Domain of Office Location = { (Lahore, Karachi, Islamabad), (Lahore, Karachi) }

- 1 • For each attribute of a relation, there is a set of permitted values, called the domain of that attribute.
- 2 • For all relations, domains of all attributes must be atomic.
- 3 • A domain is atomic if elements of the domain are considered to be indivisible units.
- 4 • The null value is a special value that signifies that the value is unknown or does not exist.

## Properties of Attributes: Example

- Suppose that we add an attribute “phone\_number” to the *instructor* relation.

- Is this attribute atomic?

Yes, bcz each no is  
an indivisible unit  
and can't be broken  
further

- Can the attribute have a null value? Perhaps yes.

bcz the instructor must  
have recently joined the uni.  
It is situation dependent

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>	<i>phone_number</i>
10101	Srinivasan	Comp. Sci.	65000	021-134567
12121	Wu	Finance	90000	042-234567
15151	Mozart	Music	40000	
22222	Einstein	Physics	95000	
32343	El Said	History	60000	
33456	Gold	Physics	87000	
45565	Katz	Comp. Sci.	75000	
58583	Califieri	History	62000	
76543	Singh	Finance	80000	
76766	Crick	Biology	72000	
83821	Brandt	Comp. Sci.	92000	
98345	Kim	Elec. Eng.	80000	null

# Database: Schema vs Instance

- **Database Schema**
  - Logical design of the database

department (dept\_name, building, budget)  
instructor (ID, name, dept\_name, salary)

- **Database Instance**
  - A snapshot of the data in the database at a give instant in time

ID	name	dept_name	salary
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

(a) The *instructor* table

dept_name	building	budget
Comp. Sci.	Taylor	100000
Biology	Watson	90000
Elec. Eng.	Taylor	85000
Music	Packard	80000
Finance	Painter	120000
History	Painter	50000
Physics	Watson	70000

(b) The *department* table

# Relation: Schema vs Instance

- Relation Schema
  - Consists of a list of attributes and their corresponding domains

department (dept\_name, building, budget)

- Relation Instance
  - A snapshot of the tuples in a relation at a give instant in time

dept_name	building	budget
Comp. Sci.	Taylor	100000
Biology	Watson	90000
Elec. Eng.	Taylor	85000
Music	Packard	80000
Finance	Painter	120000
History	Painter	50000
Physics	Watson	70000

(b) The *department* table

- We often use the same name (such as *department*) to refer to both the schema and the instance. However, when required, we can differentiate:
  - “The *department* schema”
  - “An instance of the *department* relation”

← IMP

## Database Schema: Example

```
classroom(building, room_number, capacity)
department(dept_name, building, budget)
course(course_id, title, dept_name, credits)
instructor(ID, name, dept_name, salary)
section(course_id, sec_id, semester, year, building, room_number, time_slot_id)
teaches(ID, course_id, sec_id, semester, year)
student(ID, name, dept_name, tot_cred)
takes(ID, course_id, sec_id, semester, year, grade)
advisor(s_ID, i_ID)
time_slot(time_slot_id, day, start_time, end_time)
prereq(course_id, prereq_id)
```

---

**Figure 2.8** Schema of the university database.

# Example of a Query

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

(a) The *instructor* table

<i>dept_name</i>	<i>building</i>	<i>budget</i>
Comp. Sci.	Taylor	100000
Biology	Watson	90000
Elec. Eng.	Taylor	85000
Music	Packard	80000
Finance	Painter	120000
History	Painter	50000
Physics	Watson	70000

(b) The *department* table

- Find the names of all instructors who work in Watson building

Einstein, Crick, Gold

# Example of a Query: Role of Link between Relations

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

(a) The *instructor* table

<i>dept_name</i>	<i>building</i>	<i>budget</i>
Comp. Sci.	Taylor	100000
Biology	Watson	90000
Elec. Eng.	Taylor	85000
Music	Packard	80000
Finance	Painter	120000
History	Painter	50000
Physics	Watson	70000

(b) The *department* table

- Find the names of all instructors who work in Watson building
  - Consider the department schema and the instructor schema
  - Observe that the attribute ***dept\_name*** is duplicated in both.
  - This duplication is useful in answering queries which involve multiple relations.

## Constraints on Contents/Tuples in a Relation: Examples

- Consider the following version of the department relation:

dept-name	program-director	budget
CS	Dr. Abdul Samad	100000
ECE	Dr. Farhan	85000
ISciM	Dr. Aeyaz	80000
CS	Dr. Abdul Samad	100000

## Constraints on Contents/Tuples in a Relation: Examples

- Consider the following versions of the **department** relation:

dept-name	program-director	budget
CS	Dr. Abdul Samad	100000
ECE	Dr. Farhan	85000
ISciM	Dr. Aeyaz	80000
CS	Dr. Abdul Samad	100000

1

- No two tuples in a relation are allowed to have exactly the same value for all attributes.

# Constraints on Contents/Tuples in a Relation: Examples

- Consider the following version of the **department** relation:

dept-name	program-director	budget
CS	Dr. Abdul Samad	100000
ECE	Dr. Farhan	85000
ISciM	Dr. Aeyaz	80000
CS	Dr. Waqar	100000

# Constraints on Contents/Tuples in a Relation: Examples

- Consider the following version of the department relation:

dept-name	program-director	budget
CS	Dr. Abdul Samad	100000
ECE	Dr. Farhan	85000
ISciM	Dr. Aeyaz	80000
CS	Dr. Waqar	100000

2

- Values for some attributes cannot be repeated in a relation.

# Constraints on Contents/Tuples in a Relation: Examples

- Consider the following versions of the department and instructor relations:

department

dept-name	program-director	budget
CS	Dr. Abdul Samad	100000
ECE	Dr. Farhan	85000
ISciM	Dr. Aeyaz	80000

instructor

ID	name	dept-name
2033	Dr. Abdul Samad	CS
2071	Dr. Farhan	ECE
3045	Dr. Pervez	Physics
3067	Dr. Waqar Saleem	ISciM

# Constraints on Contents/Tuples in a Relation: Examples

- Consider the following relations:

*department*

dept-name	program-director	budget
CS	Dr. Abdul Samad	100000
ECE	Dr. Farhan	85000
ISciM	Dr. Aeyaz	80000

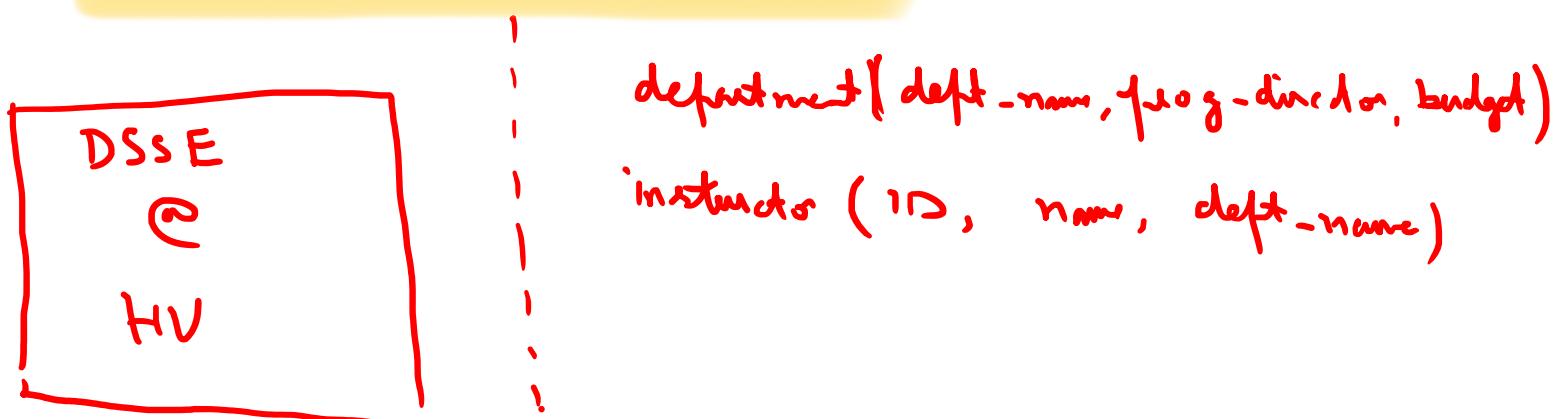
*instructor*

ID	name	dept-name
2033	Dr. Abdul Samad	CS
2071	Dr. Farhan	ECE
3045	Dr. Pervez	Physics
3067	Dr. Waqar Saleem	ISciM

- 3 • Some attributes can only be assigned values that are present in another relation.

## Constraints on Tuples in a Relation: What is the Source?

- Any relational model (that you develop) is trying to model a real-world enterprise/business.
  - For instance, in the previous slides, the relational model was an attempt at modeling the DSSE at Habib University.
- Constraints on the tuples in a relational model (such as the examples shown in previous slides) are dictated by the rules/constraints of the real-world enterprise/business being modeled.



## Constraints on Tuples in a Relation: How are these Constraints Expressed/Specified?



- Through “Keys”
- Through the following types of “Keys”, we specify how tuples within a given relation must be different from one another.
  - Superkey
  - Candidate Key
  - Primary Key
  - Foreign Key

# Keys: Superkey

- A **superkey** K is a set of one or more attributes such that no two distinct tuples can have the same values on all attributes in the set K.

Formally, let  $R$  denote the set of attributes in the schema of relation  $r$ . If we say that a subset  $K$  of  $R$  is a *superkey* for  $r$ , we are restricting consideration to instances of relations  $r$  in which no two distinct tuples have the same values on all attributes in  $K$ . That is, if  $t_1$  and  $t_2$  are in  $r$  and  $t_1 \neq t_2$ , then  $t_1.K \neq t_2.K$ .

- {ID}?  
*Super Key*
- {Name} ?  
*Not a Super Key*
- {ID, Name}?  
*Super Key*
- {name, dept\_name, salary}?

ID	name	dept_name	salary
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

(a) The *instructor* table

# Keys: Superkey

- A **superkey** K is a set of one or more attributes such that no two distinct tuples can have the same values on all attributes in the set K.

Formally, let  $R$  denote the set of attributes in the schema of relation  $r$ . If we say that a subset  $K$  of  $R$  is a *superkey* for  $r$ , we are restricting consideration to instances of relations  $r$  in which no two distinct tuples have the same values on all attributes in  $K$ . That is, if  $t_1$  and  $t_2$  are in  $r$  and  $t_1 \neq t_2$ , then  $t_1.K \neq t_2.K$ .

- {ID}?**
  - is a Superkey: enough to identify a tuple uniquely
- {Name} ?**
  - is not a Superkey: two instructor can have the same name
- {ID, Name}?**
  - is a superkey: enough to identify a tuple uniquely
- {name, dept\_name, salary}?**
  - Not a superkey. Possible for two instructors to have the same values of name, dept\_name, salary

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

(a) The *instructor* table

# Keys: Candidate Keys

- A superkey may contain extraneous attributes.
- For example, each of the following is a superkey:

- $\{ID\}$
- $\{ID, name\}$
- $\{ID, name, dept\_name\}$

- In general, if  $K$  is a superkey, then so is any superset of  $K$ .
- We are often interested in superkeys for which no proper subset is a superkey.

- Such minimal superkeys are called candidate keys.

- Only  $\{ID\}$  is the candidate key.

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

(a) The *instructor* table

## Keys: Candidate Keys

- Consider the following relation schema:

student (name, date\_of\_birth, major, student\_id, address)

- Candidate Keys?

{student\_id}

## Keys: Candidate Keys

- Consider the following relation schema:

student (name, date-of-birth, major, student-id, address)

- Candidate Keys
  - {student-id}
  - {name, address}
- It is possible that several distinct sets of attributes could serve as a candidate key.

## Keys: Primary Key

→ Have to justify

- Primary Key

- One of the candidate keys that is chosen as the principle means of identifying tuples within a relation
- The choice is made by the database designer
- Notation:

*classroom(building, room\_number, capacity)*

*department(dept\_name, building, budget)*

- Equivalent Term: Primary Key Constraint

- The designation of a key represents a constraint in the real-world enterprise being modeled through the relational model.
- Therefore, “primary keys” are also referred as “primary key constraints”.

1 CS  
 2 ~  
 ?  
 1 EE  
 2 EE  
 3 CE

## Keys as a Representation of Constraints in the Real World

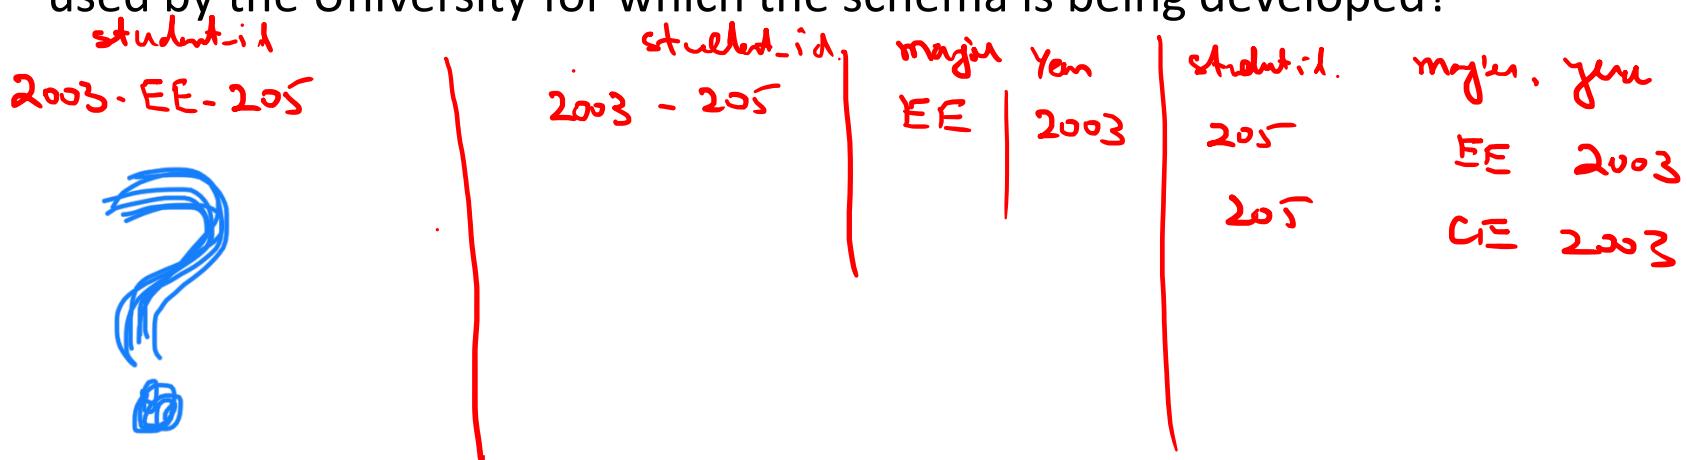
- Consider the following alternative relation schemas:

student (student-id, major, year-of-entry, date-of-birth, address)

student (student-id, major, year-of-entry, date-of-birth, address)

student (student-id, major, year-of-entry, date-of-birth, address)

- What does each option tell you about the Student ID conventions being used by the University for which the schema is being developed?



# Constraints on Contents/Tuples in a Relation: Examples

- Consider the following relations: *department*

dept-name	program-director	budget
CS	Dr. Abdul Samad	100000
ECE	Dr. Farhan	85000
ISciM	Dr. Aeyaz	80000

?

## *instructor*

ID	name	dept-name
2033	Dr. Abdul Samad	CS
2071	Dr. Farhan	ECE
3045	Dr. Pervez	Physics
3067	Dr. Waqar Saleem	ISciM

- Some attributes can only be assigned values that are present in another relation.

## Keys: Foreign Key

- Consider the following relations:

<i>department</i>		
<i>dept-name</i>	<i>program-director</i>	<i>budget</i>
CS	Dr. Abdul Samad	100000
ECE	Dr. Farhan	85000
ISciM	Dr. Aeyaz	80000

<i>instructor</i>		
<i>ID</i>	<i>name</i>	<i>dept-name</i>
2033	Dr. Abdul Samad	CS
2071	Dr. Farhan	ECE
3045	Dr. Pervez	Physics
3067	Dr. Waqar Saleem	ISciM

- Attribute *dept\_name* in relation *instructor* can only be assigned values that are present as the values of attribute *dept\_name* in relation *department*.
- dept\_name*** is a **foreign key** from the relation ***instructor*** to the relation ***department***
- Referencing relation:** *instructor*   **Referenced relation:** *department*
- The referenced attributes (*dept\_name*) must be the primary key in the referenced relation.

# Keys: Foreign Key

- Formally:

A **foreign-key constraint** from attribute(s)  $A$  of relation  $r_1$  to the primary-key  $B$  of relation  $r_2$  states that on any database instance, the value of  $A$  for each tuple in  $r_1$  must also be the value of  $B$  for some tuple in  $r_2$ . Attribute set  $A$  is called a **foreign key** from  $r_1$ , referencing  $r_2$ . The relation  $r_1$  is also called the **referencing relation** of the foreign-key constraint, and  $r_2$  is called the **referenced relation**.

department (dept-name, program-director, budget)  
Instructor(ID, name, dept-name)

<u>department</u>		
dept-name	program-director	budget
CS	Dr. Abdul Samad	100000
ECE	Dr. Farhan	85000
ISciM	Dr. Aeyaz	80000

<u>instructor</u>		
ID	name	dept-name
2033	Dr. Abdul Samad	CS
2071	Dr. Farhan	ECE
3045	Dr. Pervez	Physics
3067	Dr. Waqar Saleem	ISciM

## Keys: Exercise

Item (Item Name, Vendor Name, Price)

Customer (Customer Code, Customer Name, Address)

Vendor (Vendor Name, Contact Person, Address)

Order(Order No, Customer Code, Item Name, Vendor Name, Order Date)

- Identify Primary Keys?

{Item Name, Vendor Name}

{Customer Code}

{Vendor Name}

{Order No}

## Keys: Exercise

Item (Item Name, Vendor Name, Price)  
Customer (Customer Code, Customer Name, Address)  
Vendor (Vendor Name, Contact Person, Address)  
Order(Order No, Customer Code, Item Name, Vendor Name, Order Date)

- Identify Foreign Keys?

{ Item Name, Vendor Name }      from Order to Item

{ Vendor Name }      from Item to Vendor

{ Customer Code }      from Order to Customer

# Keys: Exercise

Item (Item Name, Vendor Name, Price)

Customer (Customer Code, Customer Name, Address)

Vendor (Vendor Name, Contact Person, Address)

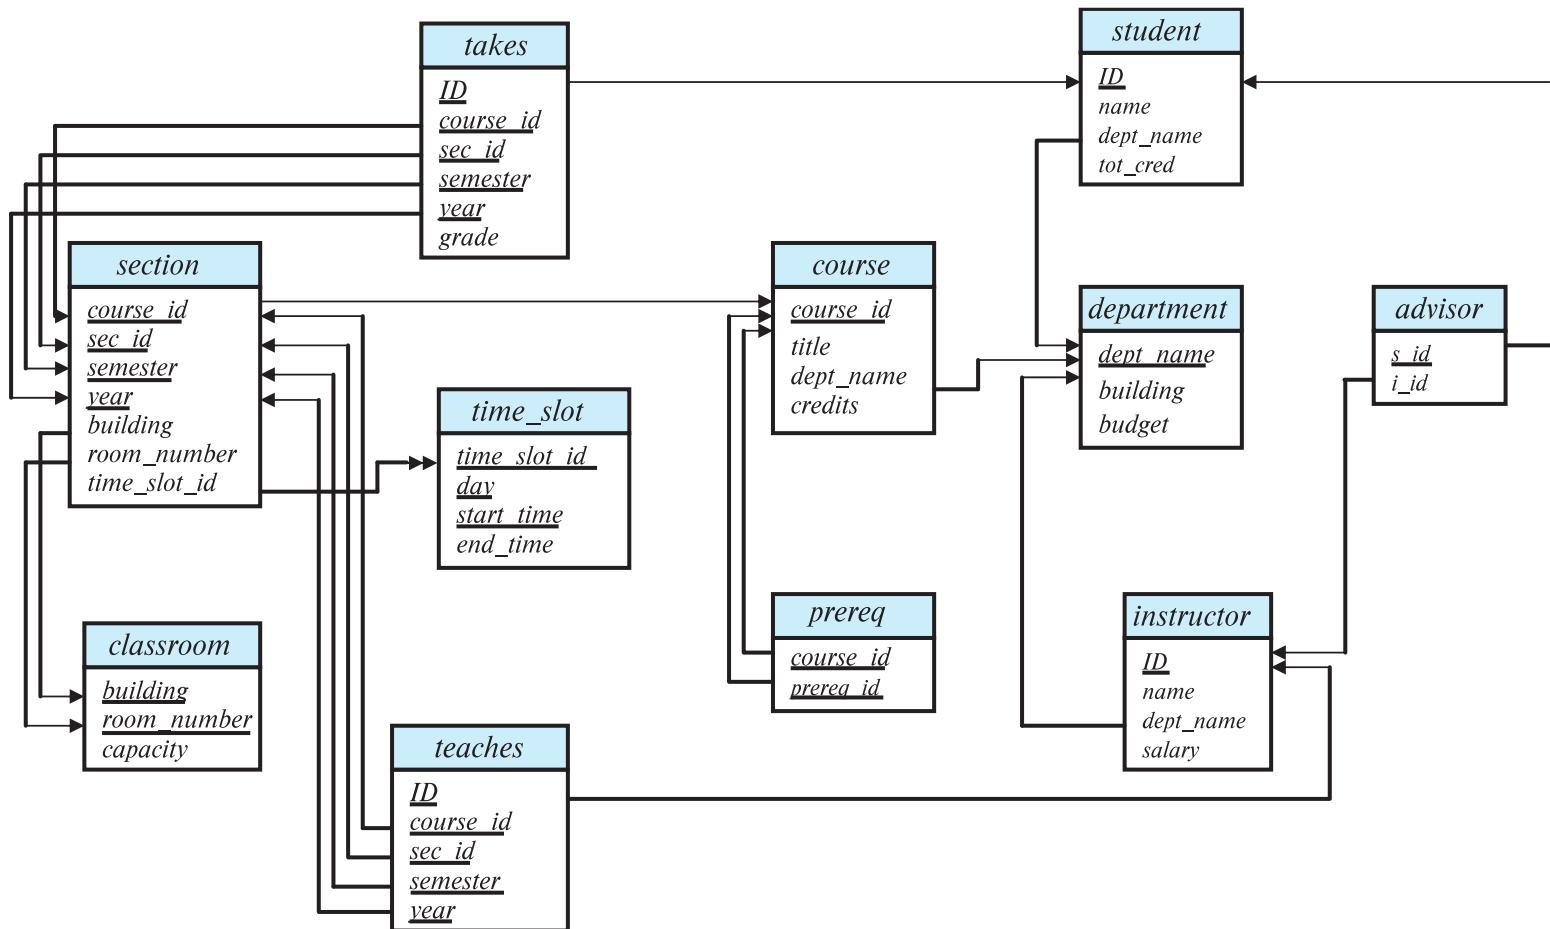
Order(Order No, Customer Code, Item Name, Vendor Name, Order Date)

- Identify Foreign Keys?

- {Customer Code} is a foreign key from Order to Customer
- {Item Name, Vendor Name} is a foreign key from Order to Item
- {Vendor Name} is a foreign key from Item to Vendor

# Schema Diagram

- Used to depict a database schema along with primary key and foreign key dependencies



# Schema Design 101

- Identify Tables
- Identify Columns/Attributes associated with each table
- Identify Primary Keys
- Identify relationships among tables through Foreign Keys

## Practice Exercise 1

- For the following relational schema
  - Identify all primary keys (underline)
  - Identify all foreign keys (use arrows)

*employee (person\_name, street, city)*  
*works (person\_name, company\_name, salary)*  
*company (company\_name, city)*

employee (CNIC, person-name, street, city)  
works (CNIC, companyID, salary)  
company (companyID, company-name, city)

## Practice Exercise 2

- For the following relational schema
  - Identify all primary keys (underline)
  - Identify all foreign keys (use arrows)

*branch(branch\_name, branch\_city, assets)*  
*customer (ID, customer\_name, customer\_street, customer\_city)*  
*loan (loan\_number, branch\_name, amount)*  
*borrower (ID, loan\_number)*  
*account (account\_number, branch\_name, balance)*  
*depositor (ID, account\_number)*

Foreign Key: { ID }

## Practice Exercise 3

- You have been asked to create a database schema for a hospital. For your help, the following relations have been identified in the schema:
  - physician
  - department
  - affiliated\_with
  - procedure
  - prescribes
  - room
  - stay
  - undergoes
  - trained\_in
  - patient
  - appointment
  - medication
- Complete the process of developing the schema